


# Neptune: A Safe, Parallel Garbage Collector for Julia

Mehmet Emre & Michael Christensen

UCSB CS263 Runtime Systems - In-Class Presentation  
May 31, 2017

A decorative light blue triangle is located in the bottom right corner of the slide.

Objective: Parallel GC in Julia

# Julia


- A scientific programming language
  - Aims for **end-to-end performance** more than latency
- Dynamically typed
- Function-level JIT compilation via LLVM
- Uses stable numeric computation libraries written in FORTRAN

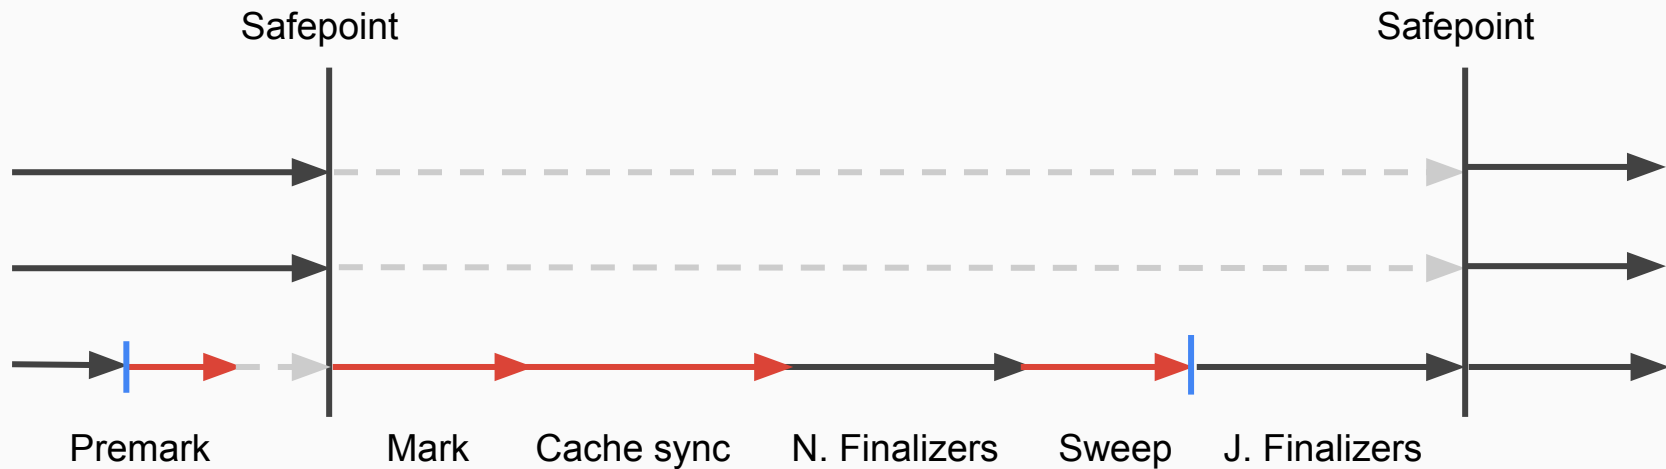
# Julia's GC (i.e. the **problem** to solve)

- Mark & Sweep
  - no copying/compacting
- Very little concurrency
  - Mostly single-threaded
  - Stop-the-world
  - Some pre-mark work concurrently
- Generational
  - 2 generations
  - Promotion: when a reachable object survives more than PROMOTE\_AGE collections
  - “remembered sets”
- Quick vs. Full
- Small objects allocated from thread-local object pools
  - Categorized by object size
- Big objects are allocated with `malloc`
- No maximum heap size
  - A huge heap (8GB) is requested from OS
  - Relies on overcommitting by OS
- GC kicks in with dynamic periods rather than internal OOM limits

# Page Manager

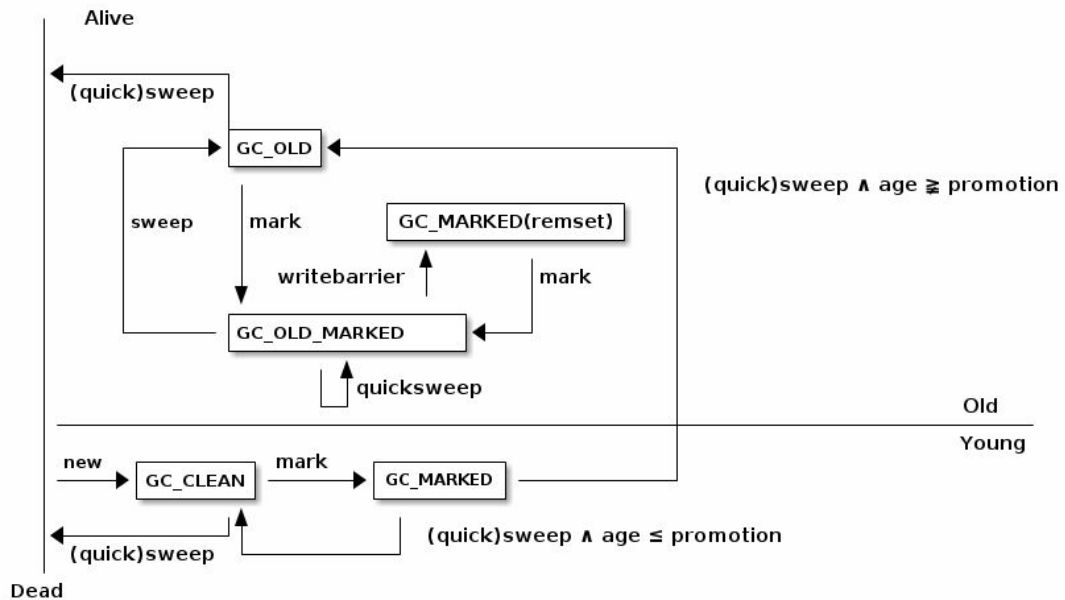
- Julia uses its own page manager for memory allocation
- The page manager organizes memory in **regions** and **pages**
- Pages in a region allocated in Rust via `mmap()` system call
- Track which chunks of pages are allocated via allocation map

[illegible]



- Doing GC
- Running Julia/native code
- → Waiting

# Object Lifetime



Neptune



# Neptune

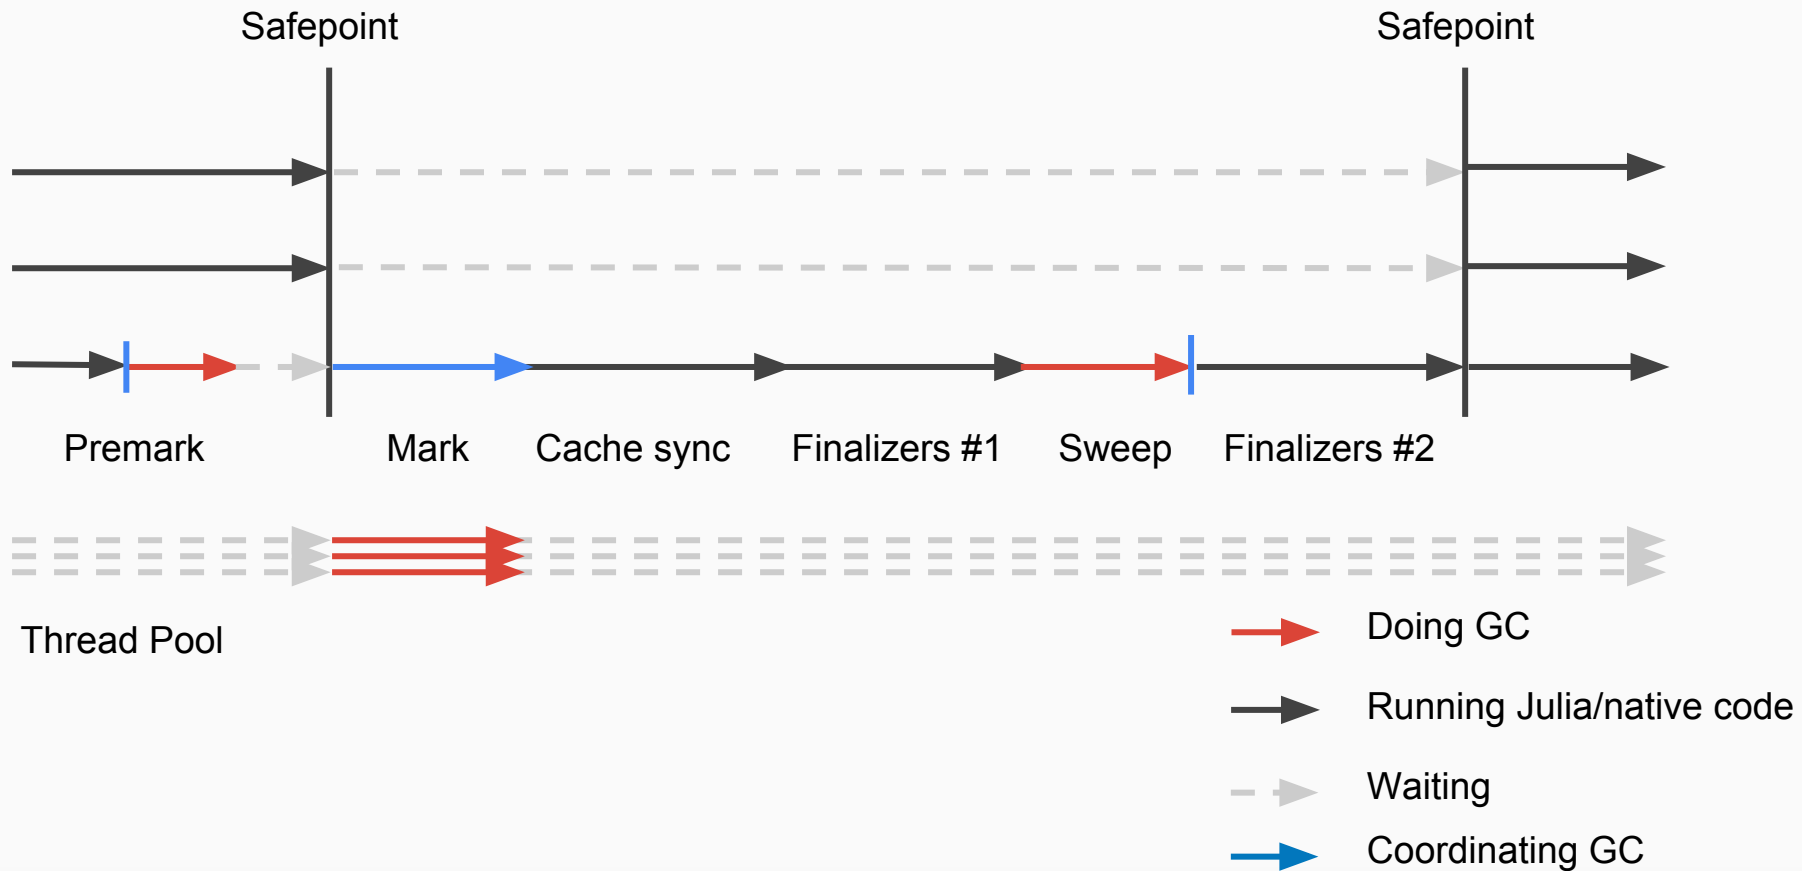
- Our parallel garbage collector
- A fixed number of threads for marking
- Single-threaded sweeping
- Mostly same as Julia's GC, to test effect of parallel GC
  - Significant effort in trying to exactly mark and sweep the same objects that Julia does, even without parallelism yet involved

# Why Rust

- A **modern** systems programming language
- Affine types for “zero-cost” abstractions
- Easy concurrency
- Memory-safety without GC
- Nice foreign function interface

# Design Decisions

- Profile-guided. Using Valgrind, OProfile, in-code timers
- Parallel sweep hindered performance
  - Sweep is mostly memory-bound, parallelizing it increased cache misses
- Heavy use of lock-free data structures and caches in many places to prevent blocking



# Thread-local mark cache

- **Remsets**
- **Big object lists**
- Statistics for decisions

# Mark stack & load balancing

- Work-stealing thread pool
- A Treiber-stack with epoch-based memory management
- Each thread
  1. Takes a job from mark stack
  2. Follows pointers until a certain depth
  3. Puts remaining pointers to the mark stack

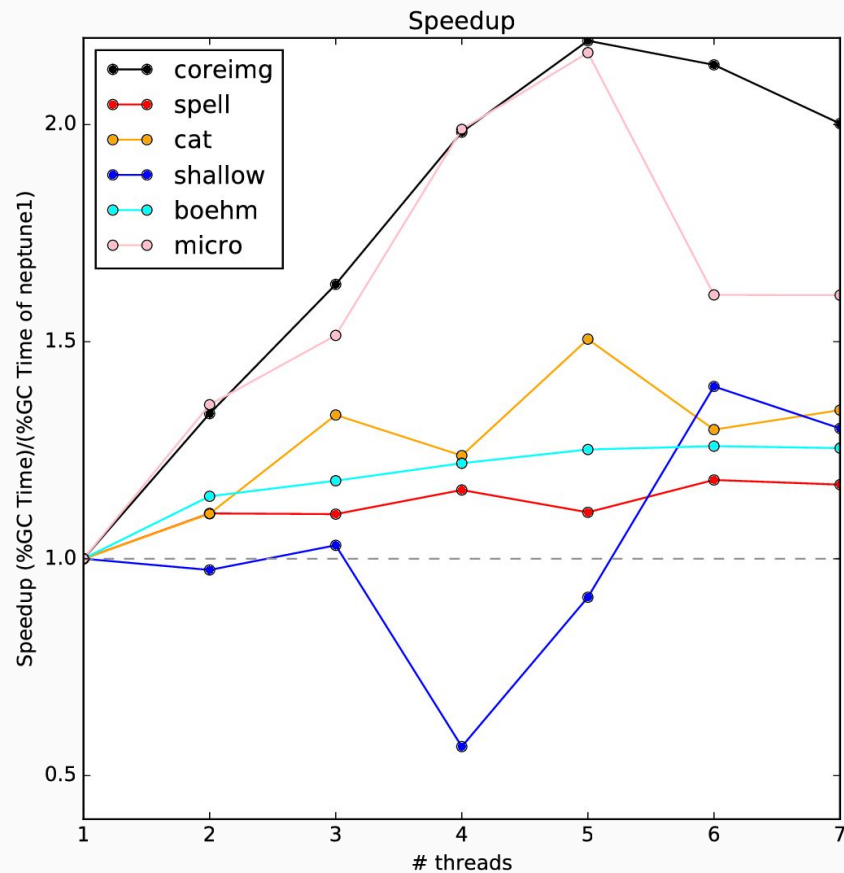

# Evaluation

# Benchmarks

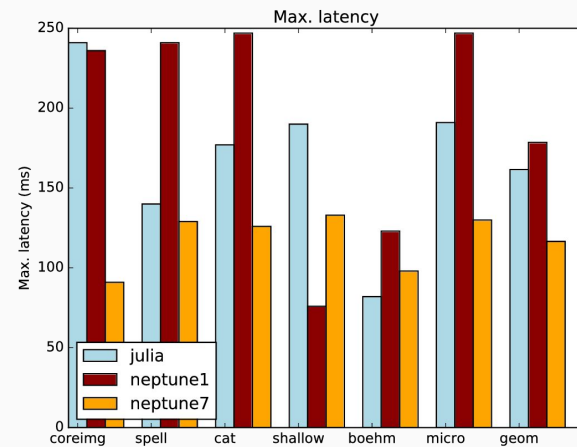

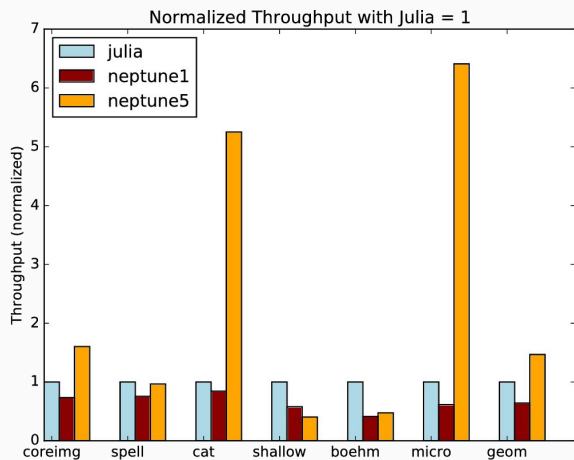
- Julia system core image (**coreimg**)
- Julia microbenchmarks (**micro**)
- Matrix concatenation (**cat**)
- Peter Norvig's spell checker (**spell**)
- Handwritten benchmark generating lots of young garbage (**shallow**)
- Hans Boehm's “An Artificial Garbage Collection Benchmark” (**boehm**)



# Speedup (in terms of %GC Time)



# Latency and Throughput



# Takeaways

# Takeaways

- Just parallelizing a GC can gain some sufficient performance
- Multi-language code performance suffers from passing the FFI boundary
- It is possible to write a GC with static thread-safety guarantees
- Profile before making performance-related decisions

# Parallel GC for Julia

It's possible and can improve performance.



“There are only two hard things in  
Computer Science: cache invalidation  
and naming things.”

---

- Phil Karlton

“There are only two hard things in  
Computer Science: cache invalidation,  
naming things and off-by-one errors.”

---

- Leon Bambrick

Don't write garbage  
collectors.

Just kidding.



Thanks!



On a memory allocation, check if GC condition is met; if so, begin the following collection steps:

**Collect:**

1. Wait for all threads to suspend
2. **Mark:**
  - a. Pointer-free objects (empty tuple type, true, false, etc.)
  - b. Each thread's remset objects (old objects pointing to younger ones, from last GC)
  - c. Each thread's local roots (thread's module, tasks, etc.)
  - d. Initial root set (builtin values, assorted constants, etc.)
  - e. Objects referenced by objects on "finalize list"
  - f. Objects put on mark stack during a. - e. above (queued to marking depth limit)
3. Flush mark cache (move big objects in cache to corresponding big object lists so sweep can see it)
4. **Sweep:**
  - a. Finalizer list (and schedule finalization)
  - b. Weak references
  - c. Malloced arrays
  - d. Big objects
  - e. Each pool in each thread
5. Run finalizers

### **Marking** Process:

1. Check current recursion depth; queue object if too high
2. Case on object's type tag, marking and recursively scanning its children/array contents/fields/etc.:
  - a. Vector
  - b. Array
  - c. Module
  - d. Task
  - e. Other data type

### **Sweeping** Process:

(Basically, depending on the object, free it...)